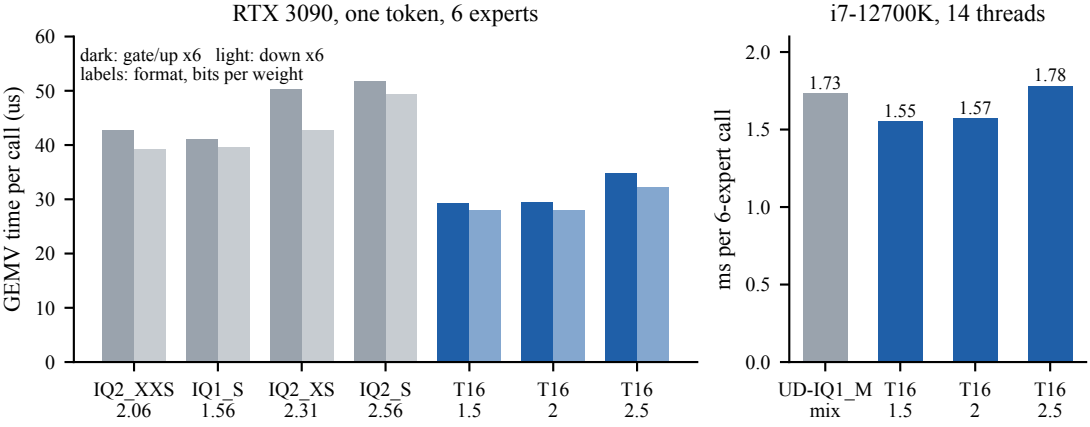
RTX 3090, one token, 6 experts
dark: gate/up x6 light: down x6
labels: format, bits per weight
GEMV time per call (us)
60
50
40
30
20
10
0
IQ2_XXS 2.06
IQ1_S 1.56
IQ2_XS 2.31
IQ2_S 2.56
T16 1.5
T16 2
T16 2.5
i7-12700K, 14 threads
ms per 6-expert call
2.0
1.5
1.0
0.5
0.0
1.73
1.55
1.57
1.78
UD-IQ1_M mix
T16 1.5
T16 2
T16 2.5