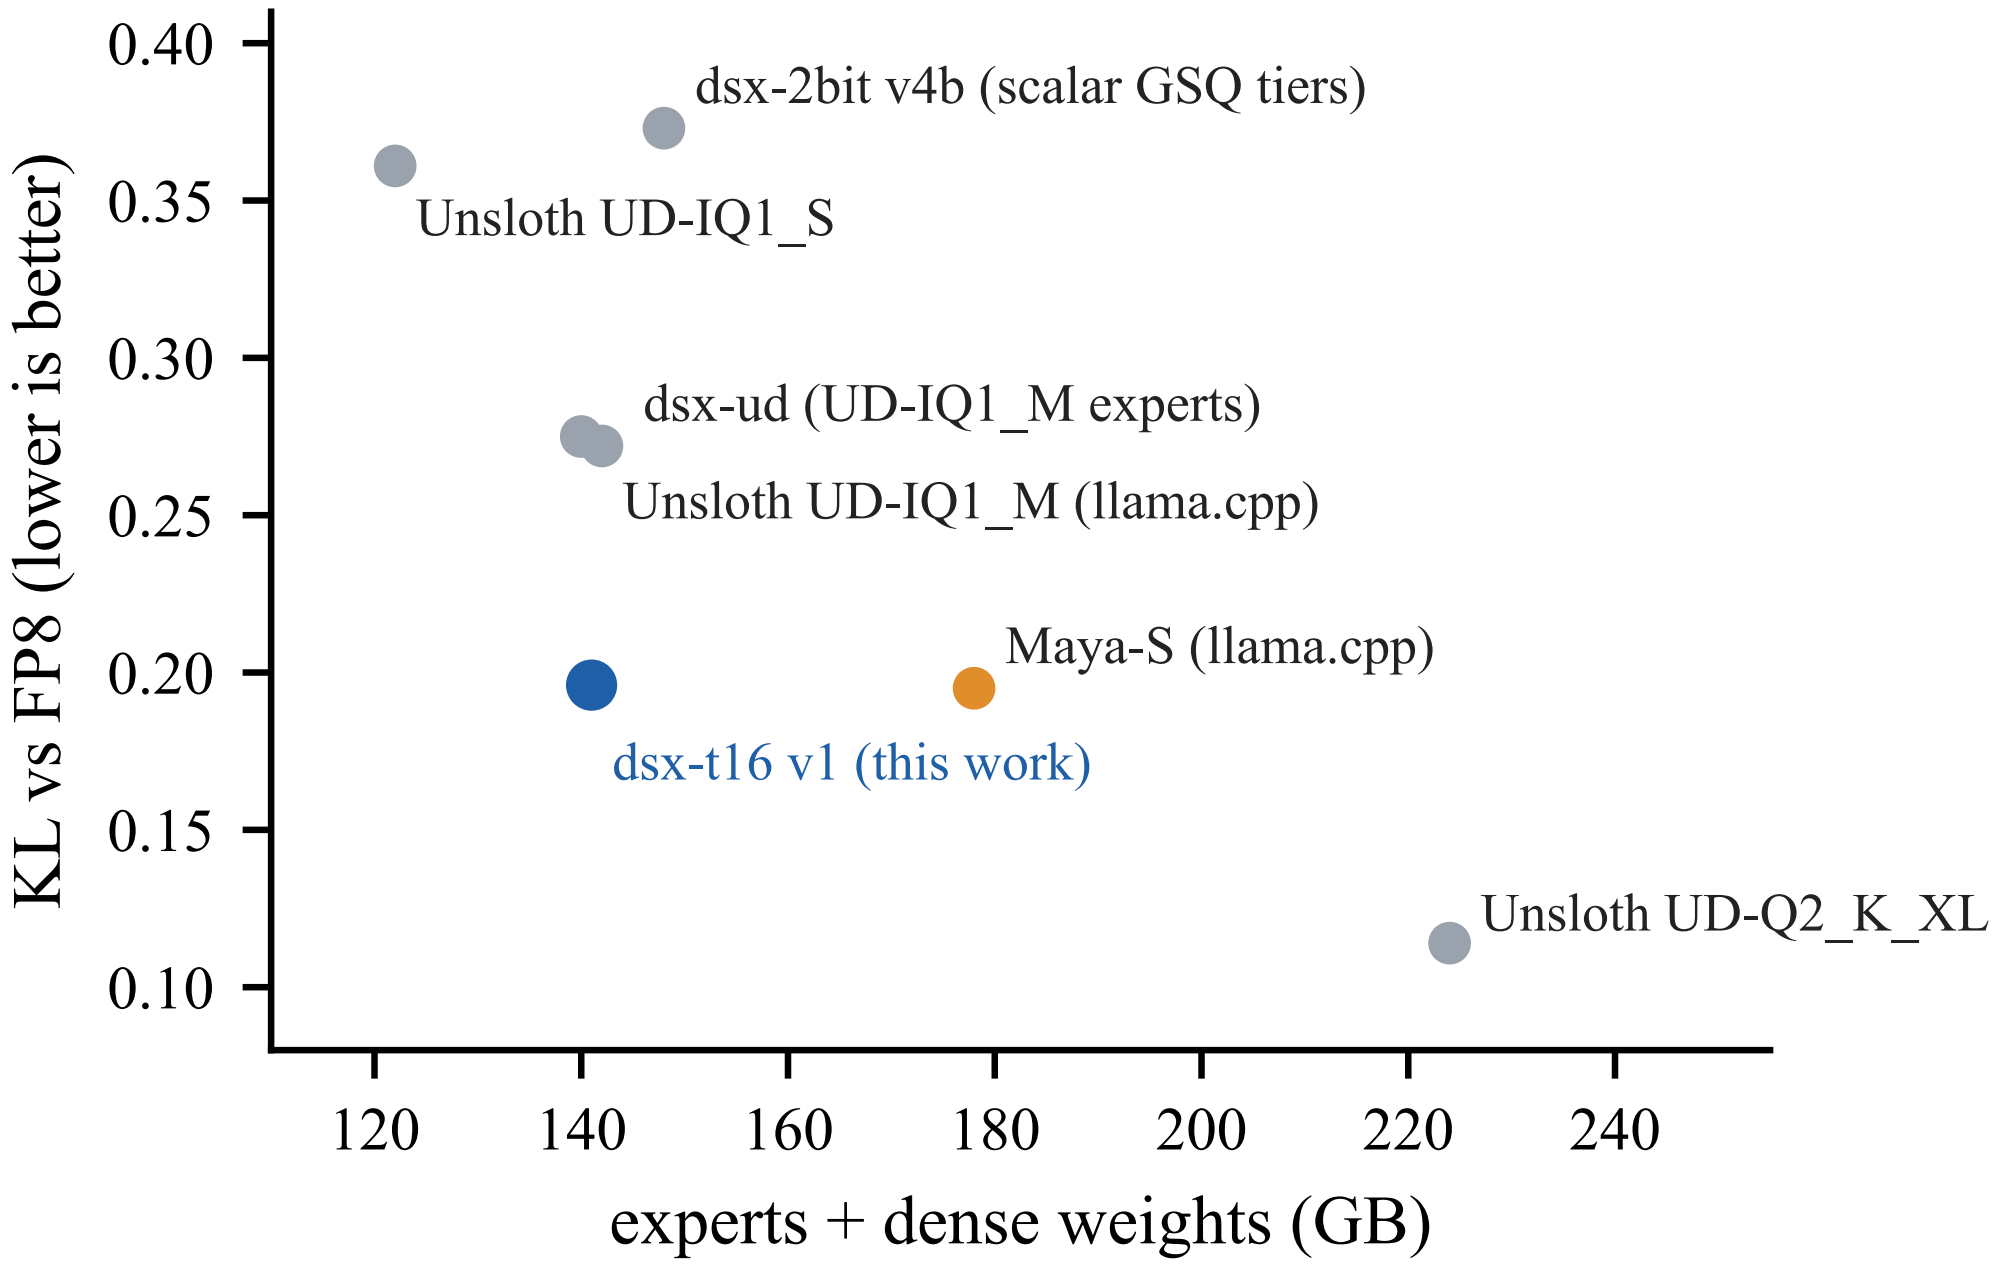
0.40
0.35
0.30
0.25
0.20
0.15
0.10
KL vs FP8 (lower is better)
dsx-2bit v4b (scalar GSQ tiers)
Unsloth UD-IQ1_S
dsx-ud (UD-IQ1_M experts)
Unsloth UD-IQ1_M (llama.cpp)
Maya-S (llama.cpp)
dsx-t16 v1 (this work)
Unsloth UD-Q2_K_XL
120
140
160
180
200
220
240
experts + dense weights (GB)